

Compilers/Linkers/Libraries

Johan Hellsvik

Example code

The lesson and the code examples are edited on

<https://github.com/PDC-support/comp-link-lib>

Obtain the code by cloning the git repository

```
git clone https://github.com/PDC-support/comp-link-lib.git
```

Contents

- Concepts: compilers, linkers, libraries
- The Gnu Compiler Collection (GCC)
- Building a Fortran library and a Fortran program with
 - static linking
 - dynamic dynamic linking
- Building and running code on large clusters
- Building a C program with dynamical libraries
 - with linking to a Cray Parallel Environment (PE) provided library
 - to a user/staff installed optional library

Concepts: compilers, linkers, libraries

- Compilers: A compiler parses the source code and generates object files and executable files. Under the hood a compiler invokes an assembler and a linker.
- Linkers: A tool to combine two or more or more object files into executable files. Linking can be static or dynamic.
- Libraries: Collection of object files for the purpose of re-use of functionality for different programs. Libraries can be static libraries or shared dynamic libraries.

Use static or dynamic linking?

- With static linking all the functionality of a program gets put in an executable file.
- Dynamic linking allows a shared library to be updated independently of the applications that are using it.
- In general dynamic linking is preferred
- Dynamic linking is the most common method on Linux systems

The GCC compiler suite

- The GNU Compiler Collection includes front ends for C, C++, Fortran and a few other languages.

Reference: <https://gcc.gnu.org/>

GCC intermediate steps

The four steps of compilation with GCC

- Preprocessing: Inserts contents of header files. Removes comments
- Compilation: Converts the preprocessed code to assembly code. The compilation is a complex process that proceed in multiple steps.
- Assembly: Converts assembly code to machine code object files
- Linking: Connects object files and libraries to executable code

Hello world in C and Fortran

C code: hello_world.c

```
#include <stdio.h>
int main() {
    printf("Hello world\n");
    return 0;
}
```

Build and run C code

```
gcc hello_world.c -o hello_world_c.x
./hello_world_c.x
```

Fortran code: hello_world.f90

```
program hello_world
    write(*,'(a)') 'Hello world'
end program hello_world
```

Build and run Fortran code

```
gfortran hello_world.f90 -o hello_world_fortran.x
./hello_world_fortran.x
```


Building a Fortran program with multiple source code files

In this example we are working with a small Fortran program with three source code files. The program is calculating the cross product of two vectors by calling a subroutine.

parameters.f90

A source code file that defines custom kinds of real numbers

```
module parameters
  implicit none
  integer, parameter :: snlprec = selected_real_kind(6, 37)    !< define precision for single reals
  integer, parameter :: dblprec = selected_real_kind(15, 307)  !< define precision for double reals
  integer, parameter :: qdprec = selected_real_kind(33, 4931) !< define precision for quad reals
end module parameters
```

crossproduct.f90

The subroutine in which the cross product is calculated

```
!cross product of two 3-vectors a and b
subroutine crossproduct(a,b,c)
  use parameters
  implicit none
  real(dblprec), dimension(3), intent(in) :: a  !< Left factor
  real(dblprec), dimension(3), intent(in) :: b  !< Right factor
  real(dblprec), dimension(3), intent(out) :: c  !< The cross product of a and b
  c(1) = a(2) * b(3) - a(3) * b(2)
  c(2) = a(3) * b(1) - a(1) * b(3)
  c(3) = a(1) * b(2) - a(2) * b(1)
  return
end subroutine crossproduct
```

mathdemo.f90

The main program. Here two vectors are defined and assigned to values, and a call made to the subroutine crossproduct.

```
program mathdemo
  use parameters
  implicit none
  real(dblprec), dimension(3) :: x,y,z
  x = (/7, 6, 3/)
  y=  (/8, 9, 5/)
  call crossproduct(x,y,z)
  write(*,10001) 'The cross product of'
  write(*,10002) 'x=', x, 'and'
  write(*,10002) 'y=', y, 'is'
  write(*,10002) 'z=', z
10001 format (a,f12.6,a)
10002 format (a,3f12.6)
end program mathdemo
```

Compiling and linking via object files

```
gfortran -c parameters.f90  
gfortran parameters.o crossproduct.f90 mathdemo.f90 -o mathdemo.x
```

```
gfortran parameters.f90 crossproduct.f90 mathdemo.f90 -o mathdemo.x
```

Building with static linking to library

Compile the source code files parameters.f90 and crossproduct.f90.

Create a library with the command `ar`

```
gfortran -c parameters.f90 crossproduct.f90  
ar r libcrossproduct.a parameters.o crossproduct.o
```

View contents of the library

```
ar t libcrossproduct.a  
parameters.o  
crossproduct.o
```

Build program mathdemo with static linking to the library

```
gfortran -static mathdemo.f90 libcrossproduct.a -o mathdemo.x
```

Building with dynamic linking to library

Compile the source code files `parameters.f90` and `crossproduct.f90`. Specify position independent code with `-fpic`. Create a library with the flag `-shared`

```
gfortran -fpic -c parameters.f90 crossproduct.f90
gfortran -shared -o libcrossproduct.so parameters.o crossproduct.o
```

Compile `mathdemo.f90`

```
gfortran -c mathdemo.f90
```

Link `mathdemo.o` to `libcrossproduct.so` and test run `mathdemo.x`

```
gfortran mathdemo.o -o mathdemo.x -L`pwd` -lcrossproduct
LD_LIBRARY_PATH=`pwd`: $LD_LIBRARY_PATH ./mathdemo.x
```

Building and running code on large clusters

- With a large pool of users comes large number of programs and libraries. This needs to be managed to avoid interference
- A common solution is to use so called module systems to selectively activate/deactivate the programs and libraries that are of relevance.
- Programs and libraries are typically built and installed with regular user privileges. Use root rights only when really needed.
- Updates of libraries need to be made with caution as they potentially can break the functionality of large number of programs

Cray programming environment (CPE)

The CPE is used on the Dardel HPE Cray EX supercomputer

Reference page: [Compilers and libraries](#)

The Cray Programming Environment (CPE) provides consistent interface to multiple compilers and libraries.

- In practice, we recommend
 - `ml cpeCray/23.12`
 - `ml cpeGNU/23.12`
 - `ml cpeAOCC/23.12`
- The `cpeCray`, `cpeGNU` and `cpeAOCC` modules are available after `ml PDC/23.12`
- No need to `module swap` or `module unload`

Compiler wrappers

- Compiler wrappers for different programming languages
 - `cc` : C compiler wrapper
 - `cc` : C++ compiler wrapper
 - `ftn` : Fortran compiler wrapper
- The wrappers choose the required compiler version and target architecture options.
- Automatically link to MPI library and math libraries
 - MPI library: `cray-mpich`
 - Math libraries: `cray-libsci` and `cray-fftw`

Compile a simple MPI code

- `hello_world_mpi.f90`

```
program hello_world_mpi
include "mpif.h"
integer myrank,mysize,ierr
call MPI_Init(ierr)
call MPI_Comm_rank(MPI_COMM_WORLD,myrank,ierr)
call MPI_Comm_size(MPI_COMM_WORLD,mysize,ierr)
write(*,*) "Processor ",myrank," of ",mysize,": Hello World!"
call MPI_Finalize(ierr)
end program
```

```
ftn hello_world_mpi.f90 -o hello_world_mpi.x
```

What flags does the `ftn` wrapper activate?

- Use the flag `-craype-verbose`

```
ftn -craype-verbose hello_world_mpi.f90 -o hello_world_mpi.x
```

Test run the hellow world MPI code

- Run on eight cores in the `shared` partition

```
user@uan01:> salloc -n 8 -t 10 -p shared -A <name-of-allocation>
user@uan01:> srun -n 8 ./hello_world_mpi.x
Processor          4  of          8 : Hello World!
Processor          6  of          8 : Hello World!
Processor          7  of          8 : Hello World!
Processor          0  of          8 : Hello World!
Processor          1  of          8 : Hello World!
Processor          2  of          8 : Hello World!
Processor          3  of          8 : Hello World!
Processor          5  of          8 : Hello World!
```

Compile a simple linear algebra code

[Link to the code](#)

Use cray-libsci

```
m1 PDC/23.12 cpeGNU/23.12
```

```
cc dgemm_test.c -o dgemm_test_craylibsci.x
```

Compile a simple linear algebra code

Use openblas

```
m1 PDC/23.12 openblas/0.3.28-cpeGNU-23.12
```

```
cc dgemm_test.c -o dgemm_test_openblas.x -I$EBROOTOPENBLAS/include -L$EBROOTOPENBLAS/lib -lopenblas
```

where the environment variable `EBROOTOPENBLAS` had been set when loading the OpenBLAS module. Its module file includes a statement

```
setenv("EBROOTOPENBLAS", "/pdc/software/23.12/eb/software/openblas/0.3.28-cpeGNU-23.12")
```

which corresponds to an export statement

```
export EBROOTOPENBLAS=/pdc/software/23.12/eb/software/openblas/0.3.28-cpeGNU-23.12
```

Check which shared libraries are needed in runtime

```
ldd dgemm_test_craylibsci.x
```

```
ldd dgemm_test_openblas.x
```

Check which shared libraries are needed in runtime

```
ldd dgemm_test_craylibsci.x
```

```
...  
libsci_cray.so.6 => /opt/cray/pe/lib64/libsci_cray.so.6  
...
```

```
ldd dgemm_test_openblas.x
```

```
...  
libopenblas.so.0 => /pdc/software/23.12/eb/software/openblas/0.3.28-cpeGNU-23.12/lib/libopenblas.so.0  
...
```


Test run the dgemm_test code

- Run on a single core in the `shared` partition

```
salloc -n 1 -t 10 -p shared -A <name-of-allocation>  
srun -n 1 ./dgemm_test_craylibsci.x  
srun -n 1 ./dgemm_test_openblas.x  
exit
```

- Expected output:

2.700	4.500	6.300	8.100	9.900	11.700	13.500
4.500	8.100	11.700	15.300	18.900	22.500	26.100
6.300	11.700	17.100	22.500	27.900	33.300	38.700

Exercise: Compile and run `fftw_test` code

```
ml cray-fftw/3.3.10.6

wget https://people.math.sc.edu/Burkardt/c_src/fftw/fftw_test.c

cc --version
cc fftw_test.c -o fftw_test.x

ldd fftw_test.x

salloc -n 1 -t 10 -p shared -A <name-of-allocation>
srun -n 1 ./fftw_test.x
```

Compilation of large program

- Examples at <https://www.pdc.kth.se/software>
- See for instance instructions for building the density functional theory program VASP https://www.pdc.kth.se/software/software/VASP/cpe23.12/6.4.3-vanilla/index_building.html

Environment variables for manual installation of software

- Environment variables for compilers

```
export CC=cc
export CXX=CC
export FC=ftn
export F77=ftn
```

- Environment variables for compiler flags
 - add `-I` , `-L` , `-l` , etc. to Makefile
- Environment variables at runtime
 - prepend to `PATH` , `LD_LIBRARY_PATH` , etc.

What happens when loading a module

```
ml show elpa/2023.05.001-cpeGNU-23.12
```

```
whatis("Description: ELPA - Eigenvalue Solvers for Petaflop-Applications")
prepend_path("CMAKE_PREFIX_PATH", "/pdc/software/23.12/eb/software/elpa/2023.05.001-cpeGNU-23.12")
prepend_path("LD_LIBRARY_PATH", "/pdc/software/23.12/eb/software/elpa/2023.05.001-cpeGNU-23.12/lib")
prepend_path("LIBRARY_PATH", "/pdc/software/23.12/eb/software/elpa/2023.05.001-cpeGNU-23.12/lib")
prepend_path("PATH", "/pdc/software/23.12/eb/software/elpa/2023.05.001-cpeGNU-23.12/bin")
prepend_path("PKG_CONFIG_PATH", "/pdc/software/23.12/eb/software/elpa/2023.05.001-cpeGNU-23.12/lib/pkgconfig")
...
```

When running your own code

- Load correct programming environment (e.g. `cpeGNU`)
- Load correct dependencies (e.g. `openblas` if your code depends on it)
- Properly prepend to environment variables (e.g. `PATH` , `LD_LIBRARY_PATH`)
- Choose correct SLURM settings

References

- [Fortran-lang.org Fortran quickstart tutorial](https://fortran-lang.org/en/latest/tutorial/)
- <https://opensource.com/article/22/5/dynamic-linking-modular-libraries-linux>
- <https://opensource.com/article/22/6/static-linking-linux>
- <https://opensource.com/article/20/6/linux-libraries>
- Compiling and running code on CPU nodes, in PDC intro presentation